

Reverse Engineering

Amedeo Cavallo

1 Introduction

Reverse engineering is a process that examines an existing product to determine detailed information and specifications in order to learn how it was made and how it works. For mechanical assemblies, this typically involves disassembly and then analyzing, measuring and documenting the parts. Reverse engineering is not limited to mechanical components or assemblies. Electronic components and computer programs (software), as well as biological, chemical and organic matter can be reverse engineered as well. [1]

The process of reverse engineering of software aims at restoring a higher-level representation (e.g. assembly code) of software in order to analyze its structure and behavior. Today, software is usually distributed in binary form which is, from an attacker's perspective, substantially harder to understand than source code. However, various techniques can be applied for analyzing binary code. [2]

2 Background

In order to figure out how to work backwards, from binary form to a higher-level representation, we need knowledge on the process that brought to that point. Assuming the reader to already have the knowledge on these processes, we will revisit some key concepts, in order to have a better understanding of the reverse engineering context.

2.1 Compilation

The compilation is the process of converting the source code into object code. It is done with the help of the compiler. The compiler checks the source code for the syntactical or structural errors, and if the source code is error-free, it then generates the object code. [3]

Considering C programming language as reference, the C compilation process is a multistage process, and can be divided into four steps, i.e., Preprocessing, Compiling, Assembling, and Linking. These steps are all typically performed automatically.

2.1.1 Preprocessor

The preprocessing passes over the source code, performing these operations:

- Comment removal
- Macro expansion
- Include expansion
- Conditional compilation (IFDEF)

2.1.2 Compiler

The code which is expanded by the preprocessor is passed to the compiler. Compiling converts the output of the preprocessor into assembly instructions.

An example of what happens when you take the classic *Hello World* program and compile it:

```
1  #include <stdio.h>
2
3  int main(int argc, char ** argv) {
4      printf("Hello!");
5      return 0
6  }
```

```
1  .LC0:
2      .string      "Hello!"
3      .text
4      .globl      main
5      .type       main, @function
6  main:
7      .LFB0:
8      .cfi_startproc
9      pushq       %rbp
10     .cfi_def_cfa_offset 16
11     .cfi_offset  6, -16
12     movq        %rsp, %rbp
13     .cfi_def_cfa_register 6
14     movl        $.LC0, %edi
15     movl        $0, %eax
16     call        printf
17     movl        $0, %eax
18     popq        %rbp
19     .cfi_def_cfa 7, 8
20     ret
21     .cfi_endproc
```

2.1.3 Assembler

The assembly code is converted into object code by using an assembler. Assemblers convert the assembly code into binary opcodes. Assuming a specific class of processors and its compatible assembly language (e.g. **x86** assembly language), i.e. the instruction `mov rax, 1` is represented by the binary opcode `0x48C7C001000000`.

2.1.4 Linker

More is needed before the object code can be executed. Mainly, all the programs written in C use library functions. These library functions are pre-compiled, and the object code of these library files is stored with `.lib` (or `.a`) extension. The main working of the linker is to combine the object code of library files with the object code of our program, i.e., if the `printf()` function is used in a program, then the linker adds its associated code in an output file. The result of linking is the final executable program.

What happens after everything has been linked is that we finally have an executable output format. The output of the compilation process can take many forms depending on the operating system:

- PE (Windows)
- ELF (Linux)
- Mach-O (OSX)
- COFF/ECOFF

This output file is often the starting point as a reverse engineer. For the scope of this document we will focus on the ELF format.

2.2 ELF Format

ELF (Executable and Linkable Format) is the object file, executable program, shared object and core file format for Linux and many UNIX operating systems. An ELF file contains an ELF header at the beginning of the file. The size of the ELF header is fixed and it contains information about the program header table and section header table. These values are zero if they are not present. [4]

The program header table, which is optional, tells how to create a process image. The section header table contains an array of `Elf32 Shdr` structures, which contain information about the various sections in the file. There can be any number of sections in the executable. Some of the common sections present in a typical binary are `.bss`, `.data`, `.dynamic`, `.debug`, `.got`, `.fini.`, `.hash`, `.interp`, `.rodata` and `.text`. [5]




Figure 1: ELF Format

The ELF files components extremely useful for reverse engineering are debug symbols. Symbols are used to aid in debugging and provide context to the loader. ELF objects contain a maximum of two symbol tables:

- **.symtab:** Symbols used for debugging / labelling
- **.dynsym:** Contains symbols needed for dynamic linking

The removal of these symbols (stripping) makes things more difficult to reverse engineer.

2.3 Stripped Binary

Stripped binaries are binaries which lack information regarding the locations, offsets, sizes and layout of functions as well as objects. Typically, all this information is stored in a symbol table which is generated by the compiler, but is removed before distribution. Although there is no performance improvement by stripping a binary it could be done for various reasons. Commercial code is stripped to make it difficult to reverse engineer proprietary algorithms; system libraries are stripped to reduce the size on disk; and malware is stripped to obfuscate it and thus complicate analysis. The general assumption across these applications is that the absence of symbol tables makes analysis of binaries more difficult. [5]

Stripped binary can be produced with the help of the compiler itself, e.g. GNU GCC compilers' -s flag, or with a dedicated tool like strip on Unix.

3 Binary Analysis

The reverse engineering process is a sequence of static and dynamic analysis that slowly refine the knowledge about the malware sample.

3.1 Static Analysis

Static analysis is performed without running the software that is to be analyzed, examining code, assets and dependencies. It is generally thought of as a more complex approach as it usually requires in-depth knowledge about the platform software is run on, frameworks being used by software and, depending on the programming language and environment used, the language software is compiled to.

Disassemblers interpret the raw bytes that represent the x86 and x64 assembly and display them in a human-readable fashion using mnemonics (e.g. translating `0xb864000000` to `mov eax, 0x40`). In case of software that is compiled down to bytecode using intermediate languages (such as Java or C#) one can use decompilers that reconstruct code that often is nearly identical to the original source-code.

3.1.1 Disassemblers

Disassemblers convert binary code to symbolic assembly language representations. A primary function of disassemblers is to visually present these assembly language representations of binary code to human analysts. To aid the analyst by organizing information, some disassemblers present control flow information or attempt to partition programs into functions. [6]

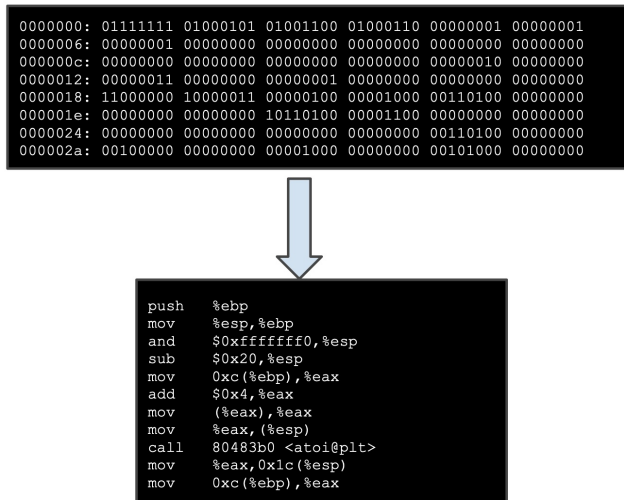


Figure 2: Static Analysis - Disassembler

Producing correct disassembly is often challenging due variable length instruction sets, data mixed with code, indirect control flow, and deliberate code obfuscation.

There are two basic techniques for disassembly [5]:

- **Linear Sweep:** This is the most straightforward and simple approach to disassembly. Examples of such a disassembler is the GNU disassembler, `objdump`. Disassembly starts from the entry point which is obtained in the header of the binary (e entry field in ELF format binaries used on most UNIX operating systems). Each successive instruction is disassembled from the next location, which is obtained by adding the length of the current instruction to the start address of the instruction. The basic disadvantage of this technique is that it cannot distinguish data from code. Any data embedded in the code is erroneously disassembled.
- **Recursive Traversal:** Recursive traversal algorithm has some advantages over linear sweep since it takes into consideration the control flow in the binary. Thus it does not misinterpret data as code. When a jump instruction is decoded, the disassembler continues disassembly from the jump target instead of blindly disassembling the next instruction. The key problem in this approach arises in the presence of indirect control flow transfer. Code that is reachable only via such transfers will not be disassembled by a vanilla recursive disassembly algorithm.

3.1.2 Decompiler

Decompilers take the process a step further and actually try to reproduce the code in a high level language. Decompilation does have its drawbacks, because lots of data and readability constructs are lost during the original compilation process, and they cannot be reproduced.

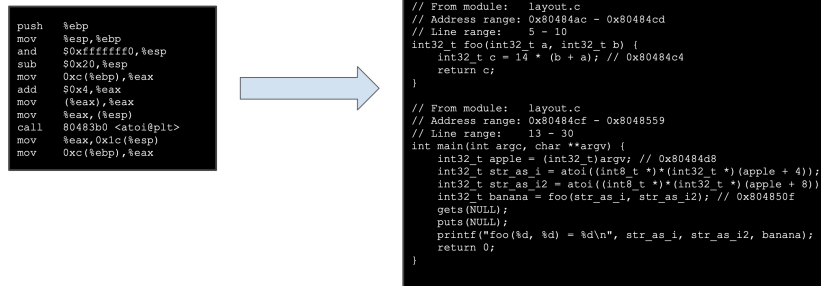


Figure 3: Static Analysis - Decompiler

3.1.3 Tools

- objdump - Disasm
- radare2 - Disasm
- capstone - Programmable Disasm
- Binary Ninja - Disasm + Primitive Decompiler
- GHIDRA - Disasm + Decompiler
- IDA Pro - Disasm + Decompiler (de facto standard)
- Angr - Binary Analysis (VEX IR) + Symbolic Execution
- rev.ng - Binary analysis with LLVM IR + Binary translation
- BAP - Binary analysis with BIL IR

3.1.4 Types

Part of the process of reverse engineering is understanding the types in the program, creating the correct structs and assigning the correct types to variables. Building correct types make decompilation more readable and easy to understand.

As an example, the same decompiled function, but on the right with the correct structs assigned to the variables:

<pre>signed int i; // [rsp+10h] [rbp-10h] signed int v2; // [rsp+14h] [rbp-Ch] void *v3; // [rsp+18h] [rbp-8h] for (i = 0; i <= 15; ++i) { if (!*(_DWORD *) (24LL * i + a1)) { printf("Size: "); v2 = sub_1AD5(); if (v2 > 0 && v2 <= 88) { v3 = calloc(v2, 1uLL); if (!v3) exit(-1); *(_DWORD *) (24LL * i + a1) = 1; *(_QWORD *) (a1 + 24LL * i + 8) = v2; *(_QWORD *) (a1 + 24LL * i + 16) = v3; printf("Chunk %d Allocated\n", (unsigned int)i); } else { puts("Invalid Size"); } return; } }</pre>	<pre>signed int i; // [rsp+10h] [rbp-10h] signed int sz; // [rsp+14h] [rbp-Ch] void *chunk; // [rsp+18h] [rbp-8h] for (i = 0; i <= 15; ++i) { if (!notes[i].state) { printf("Size: "); sz = get_long(); if (sz > 0 && sz <= 0x58) { chunk = calloc(sz, 1uLL); if (!chunk) exit(-1); notes[i].state = 1; notes[i].size = sz; notes[i].data = (__int64)chunk; printf("Chunk %d Allocated\n", (unsigned int)i); } else { puts("Invalid Size"); } return; } }</pre>
---	---

Figure 4: Static Analysis - Struct

3.2 Dynamic Analysis

Dynamic analysis is used to observe data and behaviour at runtime. It uses some very intuitive and easy to understand concepts such as memory scanners and debuggers. Memory scanners allow reverse-engineers to observe and scan memory of running programs. They can be used to find the location of values of interest and manipulate them. Debuggers instead, allow us to halt programs, execute instructions one at a time, examine registers and stack frames and trace function calls. They are especially useful to dereference pointer-chains, intercept calls to functions of interest and understand how data of interest is being accessed at runtime.

3.2.1 GDB

The GNU Debugger (GDB) is a portable debugger that runs on many Unix-like systems and works for many programming languages, including Ada, C, C++, Objective-C, Free Pascal, Fortran, Go, and partially others.

Standard GDB is not suitable to use for reverse engineering and exploit development, i.e. it still lacks a hexdump command. `pwndbg`¹ is a GDB plug-in with a focus on features needed by low-level software developers, hardware hackers, reverse engineers and exploit developers.

A small subset of useful features are:

- **Watch Expressions:** You can add expressions to be watched by the context.
- **Disassembly:** `pwndbg` uses Capstone Engine to display disassembled instructions, but also leverages its introspection into the instruction to extract memory targets and condition codes.
- **Heap Inspection:** `pwndbg` enables introspection of the glibc allocator, `ptmalloc2`, via a handful of introspection functions.
- **Process State Inspection:** Use the `procinfo` command in order to inspect the current process state, like UID, GID, Groups, SELinux context, and open file descriptors.
- **ROP Gadgets:** `pwndbg` makes using ROPGadget easy with the actual addresses in the process.
- **Finding Leaks:** Finding leak chains can be done using the `leakfind` command. It recursively inspects address ranges for pointers, and reports on all pointers found.
- **Virtual Memory Maps:** `pwndbg` enhances the standard memory map listing, and allows easy searching.

¹<https://github.com/pwndbg/pwndbg>

In order to exploit the full potential of `pwndbg`, knowledge of the most commonly used features of standard GDB are required:

- Disassemble:

```
1 set disassembly-flavor intel #sets syntax
2 disass *address #disassemble
```

- Execution:

```
1 step (s) #exec nextline - enter fun
2 next (n) #exec nextline - jump call
3 finish (f) #exec til ret
4 continue (c) #continue execution
```

- Examine:

```
1 x/numF *address #show num data of type F (bx, wx, gx, c, s)
2 printf "%c", $reg #print char from register
```

- Breakpoints:

```
1 b *address #set software breakp at addr
2 hb *address #set hardware breakp at addr
3 b *address if $reg==val #set conditional breakp
4 del br_num #remove breakpoint br_num
```

- Watchpoints:

```
1 w *address #set watch for write at addr
2 rw *address #set watch for read at addr
```

- Registers:

```
1 set $reg = val #set register to a certain value
```

- Automate:

```
1 #create command that are runned after a breakp
2 commands br_num
3     command_list
4 end
```

4 Adversarial Contest

In some applications there is a need for software developers to protect their software against reverse engineering. The protection of intellectual property (e.g. proprietary algorithms) contained in software, confidentiality reasons, and copy protection mechanisms are the most important examples. Another important aspect are cryptographic algorithms such as AES. [2] The most common example instead are malwares, where code obfuscation and analysis mitigations are applied in order to not be identified as that.

4.1 Static Analysis Mitigations

- **Complex CFG:** Create a complex control flow graph (CFG). Not aligned jumps to break the disassembler, modifying the compiler to add those instructions everywhere in the program. This usually will affect performance, so people tend to use this on code that doesn't need performance very much. Another example is adding dead code, code that is never executed.
- **Packing:** A technique to hide the real code of a program through one or more layers of compression / encryption. At run-time the unpacking routine restores the original code in memory and then executes it.
- **Header Corruption:** Most of analysis tools relies on what is written in the header of the binary, i.e. GDB and Ghidra read the header searching for symbols in the program. Messing with the header will affect the output of the decompilation process.

4.2 Dynamic Analysis Mitigations

- **Debugging Only Once:** GDB is able to debug programs calling the syscall `ptrace` with the PID of the program as a parameter. The key point is that only one instance of `ptrace` can be called simultaneously. The program can debug itself, autonomously calling `ptrace`, resulting in GDB not being able to attach to the process.
- **Check for Debugger:** `0xcc` is the byte used by debuggers for breakpoints, so continuously spamming `0xcc` to the program to execute makes debugging harder.
- **Divert Execution:** Add really complicated control flows, i.e. using multithreads programs or sending signals to trigger functions instead of calling it.

5 Walkthrough

5.1 Ghidra

Ghidra is a software reverse engineering (SRE) framework developed by NSA's Research Directorate for NSA's cybersecurity mission. It helps analyze malicious code and malware like viruses, and can give cybersecurity professionals a better understanding of potential vulnerabilities in their networks and systems. It provides a disassembler and decompiler, with a large library of supported processors / architectures.

1. Installation:

- Download the latest release from <https://ghidra-sre.org>.
- Unzip the installation bundle. This contains everything you need to run Ghidra.
- Install Java 11 64-bit Runtime and Development Kit (JDK).
- Launch Ghidra (`./ghidraRun.sh` or `./ghidraRun.bat` depending on the operating system).

2. Create Project:

- Ghidra group binaries into projects. Projects can be shared across multiple users.
- Programs and binaries can be imported into a project.
- File - New Project.
 - Select Non-Shared Project.
 - Select Directory.
 - Name the project and select Finish.

3. Load Binary:

- Import Window.
 - In this window you can inform Ghidra about the target binary.
 - Architecture / Language.
 - File format.
- Ghidra will attempt to autodetect features based on the file format. In our case these features are provided by the ELF header.
- After the file is imported, a results summary window will appear. Various file features will be listed.

5.2 revmem

A crackme (often abbreviated by cm[citation needed]) is a small program designed to test a programmer's reverse engineering skills.

5.2.1 Initial Analysis

First thing first: download the binary and try to get a basic idea of what we are dealing with:

```
acidburn@virtual:/mnt/hgfs/Shared/05$ ./revmem
Gimme the flag!
acidburn@virtual:/mnt/hgfs/Shared/05$ ./revmem flag{}
Wrong!
acidburn@virtual:/mnt/hgfs/Shared/05$ file revmem
revmem: ELF 64-bit LSB shared object, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2, for GNU/Linux 3.2.0, BuildID[sha1]=c37fc2fcdcd8b9c8c1d9a1690b4bda27dc165c4d6, stripped
acidburn@virtual:/mnt/hgfs/Shared/05$
```

Figure 5: Initial Analysis

A quick check and we see it appears to be a Linux 64 bits binary, stripped (no symbols/helpful debugging information).

5.2.2 Static Analysis

Now load Ghidra using ghidraRun, create a project then Import the revmem file (press I or drag and drop the files directly). Double click on the file, then, after using default options and enabling Analysis, you will see this workspace below.

You can notice that there is nothing listed on the right side for now:

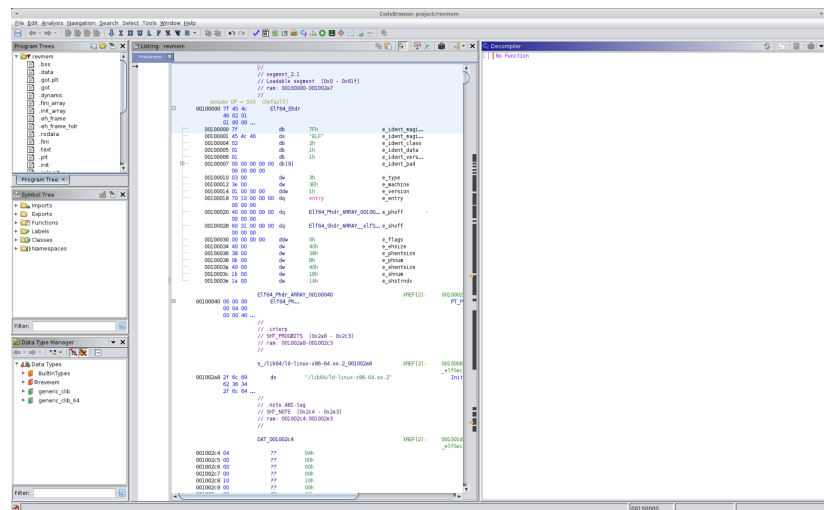


Figure 6: Workspace

Now that we are ready to start, let's hunt for the entry point and see what Ghidra can help us with.

We first want to search the main function, but as we can see quickly, there is no main function so we start looking at the entry function which is the first function called during runtime:

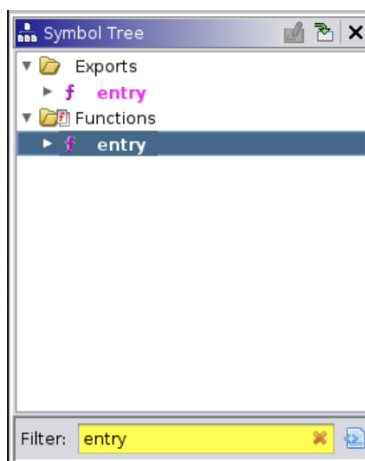


Figure 7: Symbol Tree

When selecting “entry”, notice on the right side the decompiler window gets finally busy and gives us an approximate C code, quite easy to read:

```
1 void entry(undefined8 param_1,undefined8 param_2,undefined8 param_3)
2
3
4 {
5     undefined8 in_stack_00000000;
6     undefined auStack8 [8];
7
8     __libc_start_main(FUN_001011da,in_stack_00000000,&stack0x00000008,&LAB_00101260,&DAT_001012d0,
9                     param_3,auStack8);
10    do {
11        /* WARNING: Do nothing block with infinite loop */
12    } while( true );
13 }
14
```

Figure 8: Decompiled Entry

The libc start main function shall initialize the process, call the main function with appropriate arguments, and handle the return from main.

In our case, double click on the first argument: FUN_001011da to get into our main function:

```
1
2 undefined8 FUN_001011da(int param_1,long param_2)
3
4 {
5     int iVar1;
6     char *__sl;
7
8     if (param_1 < 2) {
9         puts("Gimme the flag!");
10         /* WARNING: Subroutine does not return */
11         exit(-1);
12     }
13     __sl = (char *)FUN_00101169();
14     iVar1 = strcmp(__sl,(char **)(param_2 + 8),0x1e);
15     if (iVar1 == 0) {
16         puts("You got the flag!");
17     }
18     else {
19         puts("Wrong!");
20     }
21     return 0;
22 }
23
```

Figure 9: Main Function

Apply variable retyping and obtain a much more readable code:

```
1
2 undefined8 main(int argc,char **argv)
3
4 {
5     int iVar1;
6     char *flag;
7
8     if (argc < 2) {
9         puts("Gimme the flag!");
10         /* WARNING: Subroutine does not return */
11         exit(-1);
12     }
13     flag = (char *)compute_flag();
14     iVar1 = strcmp(flag,argv[1],0x1e);
15     if (iVar1 == 0) {
16         puts("You got the flag!");
17     }
18     else {
19         puts("Wrong!");
20     }
21     return 0;
22 }
23
```

Figure 10: Variable Type

Decompiling the `compute_flag` function and applying variable retyping will allow us to understand how the flag is being computed:

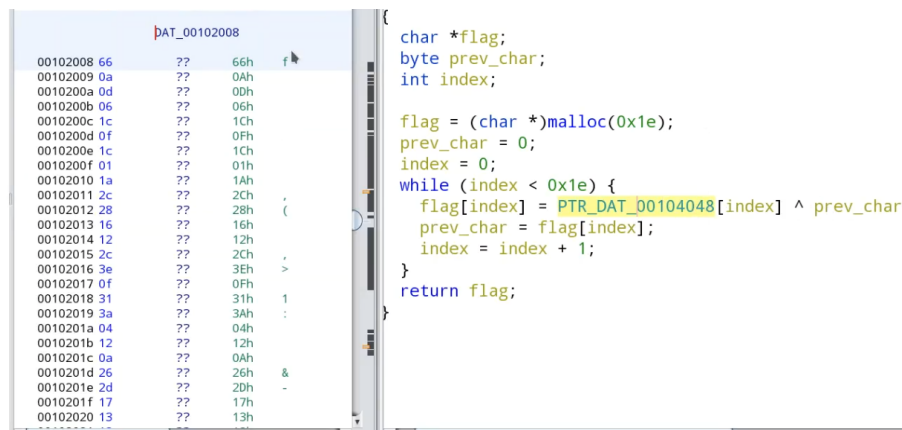


Figure 11: Compute Flag

We can simply reproduce the algorithm in a python snippet and capture the flag:

```

1 hex_str = "660A0D061C0F1C011A2C2816122C3E0F313A04120A262D171313170116186A1700
2
3 s = bytes.fromhex(hex_str)
4
5 prev_char = 0
6 index = 0
7 flag = b""
8
9 while index < 0x1e:
10     flag += bytes([s[index] ^ prev_char,])
11     prev_char = flag[-1]
12     index = index + 1

```

5.2.3 Dynamic Analysis

References

- [1] Ken Thayer. How does reverse engineering work? 2017.
- [2] Katzenbeisser S. Schrittwieser S. *Code Obfuscation against Static and Dynamic Reverse Engineering*. San Val, 2011.
- [3] Compilation process in c. 2010.
- [4] Executable and linkable format. 2010.
- [5] Arvind Ayyangar. Static disassembly of stripped binaries. Master’s thesis, Stony Brook University, 2010.
- [6] Laune Harris and Barton Miller. Practical analysis of stripped binary code. *SIGARCH Computer Architecture News*, 2005.